COMP30027 Assignment 2 Report

Anonymous

1 Introduction

The goal of this project was to predict the star ratings of a review of a restaurant based on the review text. The data available for learning and evaluation consisted of approximately 250000 Yelp reviews, and is made available by (Medhat et al., 2014) and (Rayana and Akoglu, 2015). Each review consists of the review text, a rating (1, 3 or 5 stars) and review metadata (which was not used). We sought to investigate to what extent it was possible to predict the rating of a review from the review text alone.

2 Background

This task is a type of sentiment analysis: the task of predicting sentiment from text. Performing sentiment analysis first requires extracting features from text via some encoding. One common method is to convert a text to a Bag-of-words (BOW), in which entries consist of word frequencies. This is method is simple, although has the disadvantage of producing very high dimensional features and not representing word ordering (Le and Mikolov, 2014). A more sophisticated technique is Paragraph Vector Encoding (PVE), which uses a neural network to embed texts within an abstract vector space in such a way that semantically similar texts are closer (Le and Mikolov, 2014).

Sentiment analysis tasks tend to deal with high dimension, sparse, and correlated features. This often leads to classes being linearly separable, making support vector machine (SVM) models a popular choice (Medhat et al., 2014). In terms of probabilistic classifiers, logistic regression models tend to perform well on sentiment analysis tasks since they do not assume features are independent, in contrast to other probabilistic classifiers like naïve Bayes (Medhat et al., 2014).

3 Method

We investigated models based on logistic regression (detailed in Section 3.2) and SVM (detailed in Section 3.3) techniques, in each case using the implementations made available through the Scikit-Learn Python library (Pedregosa et al., 2011). These models used PVEs of the review text, which was produced using the implementation made available through the Gensim Python library (Řehůřek and Sojka, 2010).

3.1 Feature Selection

To decide between BOW and PVE representations of review text, we used Principle Component Analysis and Singular Value Decomposition to visualise each feature in 2 dimensions and observed that PVE resulted in more distinct class clusters. Empirically, including any BOW features always decreased performance in simple models, so PVE was used exclusively for all subsequent models. This means that the dimension of the PVE-space is a hyperparameter for all of our models. This can be interpreted as a measure of complexity of the model, since as the dimension increases the PVE is better able to represent the semantic nuance of the text, and of the meaning of the text is likely irrelevant to the task-at-hand.

3.2 Logistic regression

We investigated a Multinomial Logistic Regression (MLR) model, plus the boosting and bagging homoeogenous ensemble techniques using MLR as the base classifier. MLR has one hyperparameter C, which determines how strongly the model fits to the training data and avoids misclassifying training instances.

Justifying Logistic Regression model selection

Several techniques exist for multi-class classification using Logistic Regression, such as training multiple binary Logistic Regressors (one for each class label) using a One-versus-rest scheme

and predicting based on the model with the highest score, or Ordinal Logistic Regression, which is particularly suited for ordinal class labels. One-versus-rest models are less efficient, tending to have larger standard errors (Agresti, 2002), and so were not investigated further for this reason. Ordinal Logistic Regression is not implemented as part of the Scikit-Learn Python library and thus was not investigated due to time constraints.

3.3 Support vector machine

We investigated three variants of SVM classifier: a SVM with a linear kernel and one-verses-rest multi-class classification (Linear-SVM), a SVM model with a radial basis function (RBF) kernel and one-verses-rest multi-class classification (RBF-SVM), and a SVM with a linear kernel separating the positive and negative sentiment classes, in which marginal instances are classified as having neutral sentiment (Binary-SVM).

All three SVM-based model have a regularisation coefficient C as a hyperparameter, which determines the trade-off between margin maximisation and training error minimisation. As C increases the margin is increased, however more training instances are allowed to be misclassified. The RBF-SVM also has the kernel hyperparameter γ[1]. The Binary-SVM has the probability threshold hyperparameter p_T, which is the probability a negative of positive sentiment prediction must exceed to be classified as that class.

Justifying SVM model selection

In contrast to BOW, PVE is well suited to classification by an SVM model since paragraph vectors have meaningful geometric relationships to one another (Le and Mikolov, 2014), and SVM models attempt to exploit the geometry of the feature space by fitting a separating hyperplane. PVE is also designed to transfer relationships in the meaning of text to linear relationships between paragraph vectors: "king" - "man" + "woman" = "queen" (Le and Mikolov, 2014). If we suppose that negative and positive sentiment reviews have exactly opposite meanings then we would expect the review text of these classes to be linearly separable with PVE. Preliminary testing of different kernel functions indeed showed that an SVM classifier with a linear kernel performed the best in every metric,

however an SVM classifier with a RBF kernel came close in performance so we continued with extended analysis of both.

Under these assumptions, it seems plausible that the neutral sentiment reviews will lie close to the hyperplane which separates the positive and negative sentiment reviews. This encourages the consideration of a binary SVM classifier which distinguishes positive and negative reviews, and marginal instances are classified as neutral. Rather than using the raw distance to the hyperplane, opted to use a probabilistic variant of SVM which estimates the probability of an instance belonging to each class based on its distance to the hyperplane (Platt, 1999). This allowed us to use a probability threshold p_T for neutral classification as a hyperparameter, rather than a distance threshold, which means the value of this threshold is independent of the specific text embedding[2]. Despite its appeal, results in the literature suggest that such a classifier will always result in worse performance as compared to using a standard multi-class approach (Koppel and Schler, 2006), however we were still interested in pursuing this approach.

3.4 Stacking model

We investigated Stacking as an ensemble technique to improve the performance of our base classifiers. We stacked all combinations of our best classifiers, Logistic Regression, RBF SVM and Linear SVM, including all three, and used an MLR for the meta-classifier. For every combination, we used the parameters found to be the best for each base classifier alone, including the PVE dimension.

4 Results

In order to identify the optimal feature space dimension for each model, we first evaluated our model (with hyperparameters untuned) on PVEs of dimensions 25 to 300 in steps of 25 encodings using an 80:20 stratified random holdout, only using the training set in each case to compute the PVE mapping. A random holdout method, rather than cross-validation, was used due to the computation expense of computing PVEs: producing cross-validation splits for many dimensions was out of the question. Using stratified splits was particularly important due to difference in proportion of the three

---

[1]In the Scikit-Learn implementation $\gamma/(\mathrm{var}(X) \cdot$ `num_features`), which was used as the value for the `gamma` parameter in the model constructor.

[2]The specific embedding of a particular text is only relevant in the context of the training corpus used to compute the embedding function (Le and Mikolov, 2014).

Figure 1: RBF-SVM Gridsearch

Figure 2: Binary-SVM Gridsearch


Figure 3: Learning curve

Figure 4: Binary-SVM Confusion Matrix

Figure 5: Linear-SVM Confusion Matrix

Figure 6: RBF-SVM Confusion Matrix

Figure 7: Logistic Regression Confusion Matrix

classes (69%, 29% and 8% positive, negative and neutral respectively). For the RBF-SVM and Binary-SVM models this dimension was 125, and for the Linear-SVM and MLR models this dimension was 150.

We then computed the PVE for a 5-fold stratified cross validation split at that dimension, again only using the training set of each split to produce the encoding. This pre-computed split was then used to tune the remaining hyperparameters. The use of cross-validation here helps avoid overfitting the training data when tuning the hyperparameters. Parameter values were initially adjusted in large increments, and then a finer full grid search was done on regions of interest. We found that the best hyperparameter values were: for Linear-SVM $C = 0.009$, for RBF-SVM $(C, \gamma) = (1.25, 0.6)$ (Figure 1), for Binary-SVM $(C, p_T) = 0.0025, 0.9)$ (Figure 2), and for MLR $C = 0.015$.

After tuning the hyperparameters, we evaluated the final model trained on PVEs of dimensions from 25 to 300 in steps of 25 to produce the learning curve in Figure 3.

We also evaluated our model on the 5-fold cross validation split at the optimal dimension, the results of which can be seen in Figures 4, 5, 6, 7.

We evaluated our stacked models with the computed PVEs using cross validation, using the best hyperparameters for each respective

Figure 8: Stacking Confusion Matrix

base classifier, and untuned parameters for the meta-classifier (which was MLR in each case). When the base classifiers had different optimal dimensions for PVE, we selected the dimension optimal to the classifier with better sole performance. Stacking models took a long time to train which made investigating the effect of different hyperparameters difficult. Stacked RBF-SVM and Linear-SVM had a higher accuracy and F1-score than all other stacked models, including all 4 models stacked together. The results for this model can be seen in 8.

5 Analysis

All models achieved better accuracy than the Zero-R baseline model, which has an accuracy of 69%, which shows that the models were to some degree using features to effectively make predictions. All models, including the stacking models, did not exceed an F-Score of 0.83 on held-out data. Since stacking does not improve model performance significantly, and model performance is very similar, this suggests that all models are drawing roughly the same decision boundary between the classes. This means the stacking model cannot exploit the strengths and minimise the weaknesses of the differences in the models, leading to approximately the same performance as its component models. Another possible reason for this is that the similar performance reflects the inherent difficultly of the classification task (with the selected features). It is likely that it is not possible to predict the rating entirely from the review text. Even for a human reader, the rating provides context by which to judge the sentiment of the review. For example the same review text may be interpreted as a genuine complement alongside a 5 star rating, but as sarcastic alongside a 1 star rating.

Despite having the best performance of the SVM-based classifiers on unseen data, it appears that the RBF-SVM tends to overfit the training data. In Figure 3 we can see that the performance of RBF-SVM on the training set is significantly higher than on the test data. This is in contrast to the other classifiers, which have very similar performance on the training and testing data. This suggests RBF-SVM is learning a decision boundary which is tailored to the training data, and so, despite performing marginally worse, Linear-SVM has better generalised from the training data. This is consistent with the assumption that the data should be linearly separable discussed in Section 3.3.

As we can see in Figures 4, 5, 6, 7 and 8 the main difference in performance between the models comes from their ability to distinguish the negative and neutral sentiment classes. All models were, to varying degrees, biased towards predicting the text to be more positive than it was, with the upper triangle of the confusion matrix containing the most incorrect predictions. One reason for this is that the data contains mostly positive and neutral examples. If the prior probability of a class is higher, it makes sense that a classifier would predict it more frequently and thus be more often incorrect when guessing this class. Another factor in this may be is that the distinction between negative and neutral, and neutral and positive sentiment language is not very clear. These two factors combined would lead to the observed tendency of models to frequently predict one class more positive than the true class.

The Binary-SVM model (F-score of 0.758) performed worse than the Linear-SVM model (F-score of 0.821) which is consistent with results in the literature (Koppel and Schler, 2006). This means that the assumptions about the text data discussed in Section 3.3 which would lead to good performance of the Binary-SVM model are too strong. It is likely that the language used in positive and negative reviews is not completely opposite, so PVEs of these reviews do not lie along a line. Even if this were the case, there is no reason an SVM model using a standard multi-class classification technique could not learn the same decision boundary as a model like Binary-SVM. However, Binary-SVM was by far the fastest to train, only needing one hyperplane and no neutral examples to build the model.

One major improvement we could make would be to attempt to simplify texts linguistically prior to producing the PVEs. In the simplest case this could involve using an automated spell checker on the reviews, and in more

complex cases undoing negatives ("not bad" ↦ "good") and identifying common phrases or idioms. Given enough examples, the paragraph vector encoding should be able to account for this without preprocessing, since words and phrases with similar meanings should be close in the feature space. However it is likely that there are not enough examples of misspellings and unusual phrasings in the training set to do this effectively in all cases.

6   Conclusion

We found that all investigated models were able to predict the rating from the review text along better than a Zero-R baseline classifier. Except for Binary-SVM, all models had very similar performance. This invites the question of whether it is possible to do better than the models discussed here with the selected features, or if that limit is due to the intrinsic predictability of the data with the features used.

This work also raises the question of to what extent linguistic preprocessing can improve performance on this dataset.

References

Alan Agresti. 2002. *Categorical data analysis. 2nd ed.* Chichester: Wiley, 2nd ed. edition.

Moshe Koppel and Jonathan Schler. 2006. The importance of neutral examples for learning sentiment. 22(2):100–109.

Quoc V. Le and Tomas Mikolov. 2014. Distributed representations of sentences and documents.

Walaa Medhat, Ahmed Hassan, and Hoda Korashy. 2014. Sentiment analysis algorithms and applications: A survey. 5(4):1093–1113.

Arjun Mukherjee, Vivek Venkataraman, Bing Liu, and Natalie Glance. 2013. What Yelp fake review filter might be doing? In *7th International AAAI Conference on Weblogs and Social Media*.

F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. 2011. Scikit-learn: Machine learning in python. 12:2825–2830.

John Platt. 1999. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. 10(3):61–74. Publisher: Cambridge, MA.

Shebuti Rayana and Leman Akoglu. 2015. Collective opinion spam detection: Bridging review networks and metadata. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 985–994.

Radim Řehůřek and Petr Sojka. 2010. Software framework for topic modelling with large corpora. In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, pages 45–50. ELRA.